Errata for Peyton, Huber and Coppock (2022)

The Generalizability of Online Experiments Conducted During the COVID-19 Pandemic

Sixteen corrections to the article and its online appendix, none of which changes a conclusion. Every corrected value below is computed at render time from the maintained reanalysis of the deposited replication archive, `10.7910/DVN/38UTBF`. The article is `10.1017/XPS.2021.17`. A corrigendum published in 2023, `10.1017/XPS.2022.23`, corrected the counts of statistically significant differences in the article's Figures 2 and 3 and in the first paragraph of page 6; entries 11 through 13 below are that corrigendum's corrections rather than this reanalysis's, carried here so the note is a complete account of what the published article gets wrong. Their values are recomputed from the reanalysis like every other, not copied out of the corrigendum, and all four agree with it exactly. The reanalysis is at `active-maintenance-aec/peyton_huber_coppock_2022`.

Values on the percentage-point scale are stated to one decimal place throughout, including numbers that were not wrong. Bold marks the tokens that changed.

---

1. The four estimates attributed to the original study are the pre-COVID replication's

Published, app. p. 18, Section A.8:

> In the original study, the 90/70 condition caused an increase in the proportion of subject that preferred the nuclear option by about **37** percentage points relative to the 90/90 condition. The estimated effect of the 90/45 condition, relative to the 90/90 condition, was **51** percentage points. Similarly, the 90/70 condition caused an increase in the proportion of subjects that approved of the nuclear option by about **17** percentage points, and the 90/45 condition caused an increase of about **27** percentage points.

Corrected:

> In the original study, the 90/70 condition caused an increase in the proportion of subject that preferred the nuclear option by about **29.7** percentage points relative to the 90/90 condition. The estimated effect of the 90/45 condition, relative to the 90/90 condition, was **47.6** percentage points. Similarly, the 90/70 condition caused an increase in the proportion of subjects that approved of the nuclear option by about **5.1** percentage points, and the 90/45 condition caused an increase of about **27.5** percentage points.

All four published figures are the estimates from the Aronow, Baron and Pinson (2019) pre-COVID replication rather than from Press, Sagan and Valentino (2013). The appendix's own Table 2 prints both columns and gives the original study 0.30, 0.48, 0.05 and 0.28 against the replication's 0.37, 0.51, 0.17 and 0.27.

2. The count of significant differences

Published, app. p. 25, Section A.10, n. 5:

> Although all 6 replication estimates are smaller in magnitude than those in the original study, only **two** of these differences (**John Podesta and Trump** stories) are statistically distinguishable from zero.

Corrected:

> Although all 6 replication estimates are smaller in magnitude than those in the original study, only **four** of these differences (**John Podesta, Trump, Scaramucci and Vermont** stories) are statistically distinguishable from zero.

The two additional differences are significant at conventional levels before any adjustment for multiple comparisons, which is what the footnote describes; three of the four survive a Benjamini-Hochberg adjustment across the six stories. The 2023 corrigendum corrected the same class of error in the article's Figures 2 and 3 and did not reach this footnote.

3. The replication estimate and its distance from the benchmark

Published, app. p. 1, Section A.1:

> Pooling across the pre-COVID studies, the summary effect size estimate is a **30** percentage point increase in support for Russian journalists. Our effect estimate of **25.5** points is 83% of this magnitude, and this difference of **5.5** percentage points is not statistically significant (P = 0.16).

Corrected:

> Pooling across the pre-COVID studies, the summary effect size estimate is a **30.5** percentage point increase in support for Russian journalists. Our effect estimate of **25.4** points is 83% of this magnitude, and this difference of **5.1** percentage points is not statistically significant (P = 0.16).

The sentence's own three figures do not agree with each other: 30 less 25.5 is 4.5, not 5.5. The ratio and the p-value it also prints are both right, and both follow from the corrected estimate and difference.

4. The non-conjoint correspondence

Published, p. 384:

> The 24 of 28 correctly signed estimates from the other replications were, on average, **49%** as large as the pre-COVID summary effect sizes.

Corrected:

> The 24 of 28 correctly signed estimates from the other replications were, on average, **48%** as large as the pre-COVID summary effect sizes.

The mean ratio over those 24 pairs is 0.483. The overall 73% and the conjoint 87% stated in the same paragraph both reproduce exactly, and 73% is the mean over the 65 correctly signed pairs of the two groups together, which is the denominator the paragraph's three figures share.

5. The average duration among browser respondents

Published, p. 388:

> Those from web applications spent approximately 7 min less time completing surveys than those from web browsers, who spent an average of **21.5** min.

Corrected:

> Those from web applications spent approximately 7 min less time completing surveys than those from web browsers, who spent an average of **21.0** min.

21.5 minutes is the average among nonmobile respondents, which the next sentence uses.

6. The inattentive share

Published, p. 389:

> If 73% of respondents were attentive and **26%** inattentive, our replication estimate would, on average, match our pre-COVID benchmarks.

Corrected:

> If 73% of respondents were attentive and **27%** inattentive, our replication estimate would, on average, match our pre-COVID benchmarks.

The correspondence is 73.0 per cent, so the two shares are the attentive figure and its complement.

7. The number of estimated AMCEs

Published, app. p. 22, Section A.9:

> In total, there are **81** estimated AMCEs in Figure A.10 – 41 for the original study and 41 for the replication.

Corrected:

> In total, there are **82** estimated AMCEs in Figure A.10 – 41 for the original study and 41 for the replication.

The two component counts are right and Figure A.10 plots 82 points.

8. The 90/70 condition's success rate

Published, app. p. 17, Section A.8:

> 2) a "90/70" condition in which the conventional strike had a **75%** chance of success

Corrected:

> 2) a "90/70" condition in which the conventional strike had a **70%** chance of success

The condition's two siblings are described at the rates their labels name, a 90% chance for the 90/90 condition and 45% for the 90/45 condition, and the label is written 90/70 in the same paragraph, in the sentence that follows it and on Figure A.8.

9. The standard error of the difference between the two intentionality arms

Published, app. p. 16, Section A.7:

> The estimated treatment effect for the COVID-specific replication is about 0.01 points smaller than the direct replication (SE = **0.04**, P = 0.46).

Corrected:

> The estimated treatment effect for the COVID-specific replication is about 0.01 points smaller than the direct replication (SE = **0.06**, P = 0.46).

0.04 is the standard error of the difference against the pre-COVID benchmark, given in the sentence that follows. The published P = 0.46 is the one-tailed p-value at the corrected standard error, so the sentence is internally consistent once the standard error is corrected.

10. The Easy ACQ pass rate among web application respondents

Published, p. 389, Table 2:

> Easy, Web app: **0.55 (0.02)**

Corrected:

> Easy, Web app: **0.54 (0.02)**

The deposited script rounds each pass rate to three decimal places before formatting it to two, so 0.5449 is printed as 0.545 and then as 0.55. At the table's own precision the rate is 0.54. It is the only one of Table 2's 36 published numbers the reanalysis does not reproduce: no other cell, difference or standard error changes, and neither does the sign or the significance of any difference in the table.

11. The count of significant differences under Figure 2

Published, p. 385, the note to Figure 2:

> Estimated summary effect sizes and 95% confidence intervals from pre-COVID experiments (horizontal axis) and COVID era replications (vertical axis). **Thirteen** of 28 replication estimates were significantly different from their pre-COVID benchmark at P $<$ **0.05.**

Corrected:

> Estimated summary effect sizes and 95% confidence intervals from pre-COVID experiments (horizontal axis) and COVID era replications (vertical axis). **14** of 28 replication estimates were significantly different from their pre-COVID benchmark at P $<$ **0.05, and 13 of these (denoted by square shaped points) remained statistically significant after adjusting for multiple comparisons to control the false discovery rate (Benjamini and Hochberg, 1995).**

The corrigendum's own account is that the code adjusting p-values for multiple comparisons was wrong, so the published note reported the adjusted count while describing it as the unadjusted one. The figure changes with the note: the corrected artwork marks the estimates that survive the adjustment with square points, which the published figure does not distinguish. No point estimate, standard error, confidence interval or substantive conclusion changes.

12. The count of significant differences under Figure 3

Published, p. 386, the note to Figure 3:

> Estimated summary effect sizes and 95% confidence intervals from pre-COVID experiments (horizontal axis) and COVID era replications (vertical axis). **Seven** of 41 replication estimates were significantly different from their pre-COVID benchmark at P $<$ **0.05.**

Corrected:

> Estimated summary effect sizes and 95% confidence intervals from pre-COVID experiments (horizontal axis) and COVID era replications (vertical axis). **12** of 41 replication estimates were significantly different from their pre-COVID benchmark at P < **0.05, and 7 of these (denoted by square shaped points) remained statistically significant after adjusting for multiple comparisons to control the false discovery rate (Benjamini and Hochberg, 1995).**

The same error as entry 11 and the same remedy. The corrigendum sets both figures' notes in the same words, and the corrected artwork again marks with square points the estimates that survive the adjustment.

13. The counts of significant differences in the first paragraph

Published, p. 384:

> Of the 24 correctly signed estimates, **10** were significantly smaller in replication. Of the four incorrectly signed estimates, three were significantly different – the foreign aid misperceptions study and two of six estimates from the atomic aversion study. Figure 3 plots the analogous information for the 41 conjoint estimates and their 41 replications, all of which are signed in the same direction. Of these, 35 of 41 were smaller in replication (**6** statistically significant differences) and 6 of 41 were larger in replication (1 of 6 significant differences).

Corrected:

> Of the 24 correctly signed estimates, **11** were significantly smaller in replication. Of the four incorrectly signed estimates, three were significantly different – the foreign aid misperceptions study and two of six estimates from the atomic aversion study. Figure 3 plots the analogous information for the 41 conjoint estimates and their 41 replications, all of which are signed in the same direction. Of these, 35 of 41 were smaller in replication (**11** statistically significant differences) and 6 of 41 were larger in replication (1 of 6 significant differences).

The corrigendum prints the whole paragraph and marks these two counts. Both are unadjusted counts that the published article took from the faulty adjustment, which is the error entries 11 and 12 correct in the two figures. The two sentences the corrigendum leaves standing are quoted here because it quotes them: they are what fixes which counts the two numbers refer to.

14. Whether every replication estimate was smaller

Published, p. 384:

> **All** replication summary estimates were smaller in magnitude than their pre-COVID estimates, with 24 of 28 signed in the same direction.

Corrected:

> **All but one of the** replication summary estimates were smaller in magnitude than their pre-COVID estimates, with 24 of 28 signed in the same direction.

The exception is the support-for-redistribution pair, whose pre-COVID summary effect is 0.00008 and whose COVID-era estimate is -0.044. It is one of the four pairs signed against their benchmark, so the count in the second half of the sentence is unaffected.

15. A reference to a table that does not exist

Published, app. p. 11, Section A.5:

> **Table A.5** provides a summary of differences between the original study and the replications by treatment arm.

Corrected:

> **Table 1** provides a summary of differences between the original study and the replications by treatment arm.

The appendix carries no Table A.5. The table this sentence points at is captioned Table 1 and appears on the following page; the deposited file that builds it is named `appendix_table_a5.tex`, so it is the printed caption rather than the reference that departed from the intended numbering.

16. A coauthor's given name set as her surname

Published, p. 384, Table 1:

> Inequality and system justification (**Trump & Ariel**, 2018)

Corrected:

> Inequality and system justification (**Trump & White**, 2018)

The article's own reference list gives the work as "Trump, Kris-Stella and Ariel White. 2018", and the body cites it correctly as Trump and White (2018). Only Table 1's row for the study carries the wrong surname. The authors' accepted manuscript types "(Trump & White, 2018)", so the substitution was made after acceptance. Neither the reproduction nor an audit of the reference list could have found it: the table prints no estimate, and the reference it mis-cites is itself correct.